

Vamvas Ioannis A.M.: 2943

Gewrgoulas Vasilis A.M.: 2954

### **Step 1**

Originally in the `handle_vfs_reply()` function of the archive `usr/src/servers/pm/main.c`, we can see `PM_FORK_REPLY` field that's where the execution of the main `fork()` happens in `start_user()` of `/usr/src/servers/pm/schedule.c`. There it's the first change we're going to make to get the group through a process in the `do_start_scheduling()` of the archive `usr/src/servers/sched/schedule.c`. In `start_user()`, we have the function of `start_inherit()`, which accepted 5 arguments. We changed her prototype through the archive `usr/src/include/minix/sched.h` so that it now accepts 6 arguments where the last argument will be the process group (`rmp->procgrp`). The function of the `start_inherit()` is then implemented in the `usr/src/lib/libsys/sched_start.c` where we put the 6th argument in the field `m9_15` of the message that will be posted to the sched (which will be `SCHEDULING_INHERIT`). We put it on the process. in the field of `m9_15` because we concluded that in the `SCHEDULING_INHERIT` don't use this field of message (more specifically we saw it through the archive `usr/src/include/minix/com.h`). Then we see that with the message `SCHEDULING_INHERIT` it'll end up in the archive `usr/src/servers/sched/main.c` where from there we see that for the case `SCHEDULING_INHERIT` it'll end up with the function `do_start_scheduling()` mentioned above, where there will be the initialization of the team of each process by taking the field `m9_15` of the message.

### **Step 2**

In the archive `usr/src/servers/sched/schedproc.h` we added the 4 fields requested in the exercise, more specifically the `pid_t procgrp` (where is the team driver), the unsigned `proc_usage` which is the use

process, the unsigned `grp_usage` which is the use of the process group and the unsigned `fss_priority` which is the priority based on the algorithm of fair routing. These 4 fields, the initials in the `do_start_scheduling()` function of `usr/src/servers/schedul/schedule.c` and more specifically, we initialize the `procgrp` field as `m_ptr->m9_15` right after the introduction of the process into the `NG_INHERIT` structure.

The field `proc_usage` initialize to 0 because the use of the process once it enters `do_start_scheduling` is 0 (because the "programming" of the process from the core has not yet beendone).

The field `grp_usage` initialize as much as the `grp_usage` of the <sup>1st</sup> process (parent) because we know that all processes of the same group should have the same `grp_usage`.

The field `fss_priority` be initialized on the basis of the type, where the `number_of_groups` find it through the function `num_of_grps()`, which is returned to `number_of_groups`.  
`num_of_grps()` we first find all processes useful and we kept in 1 table the drivers of the process group and in the We've always counted the different group drivers we find. within this table where it will finally be and the `number_of_groups` we're looking for.

Then the information of the 4 above fields is done in the `do_noquantum()` function of the `usr/src/servers/schedul/schedule.c` file, where we check if we have user processes and if we then have the following information:

The field `proc_usage` process that finished its quantum, we increase it by `RMP->time_slice` if `time_slice` is equal to `USER_QUANTUM` otherwise we increase it by `USER_QUANTUM`.

Then we run all the processes we have in the structure of `schedproc.h` and increase all `grp_usage` of the processes that have the same team driver as the process that ended its quantum `USER_QUANTUM`.

Finally, for all user processes, we update their data as follows:

The  $\text{proc\_usage} = \text{proc\_usage} / 2$

The  $\text{grp\_usage} = \text{grp\_usage} / 2$

And the  $\text{fss\_priority}$  based on the type  $\text{fss\_priority} = \text{proc\_usage} / 2 + \text{grp\_usage} * \text{number\_of\_groups} / 4 + \text{base}$ , where  $\text{base} = 0$  (is the same type that was climbed above step 2 in terms of the initialization of the  $\text{fss\_priority}$ ).

### **Step 3**

To make the processes useful in 1 queue only (as you ask in the pronunciation), we changed the file `usr/src/include/minix/config.h` and now we put as  $\text{NR\_SCHED\_QUEUES} = 8$  and  $\text{MAX\_USER\_Q} = \text{MIN\_USER\_Q} = \text{USER\_Q} = 7$  so that the queues 0-6 are the same as before,  $\text{MAX\_USER\_Q} - \text{MIN\_USER\_Q} + 1$  be the tail useful and the tail 8 be the tail idle (as it was in the previous version of the minix tails). Then, in order to pass the  $\text{fss\_priority}$  to the core, we take the six steps :

- 1) We change the prototype of the `sys_schedule` so that instead for 4 argument, accept 5, so that its 5<sup>th</sup> field is  $\text{fss\_priority}$ . This was done through the `usr/src/include/minix/syslib.h` file.
- 2) In the `usr/src/lib/libsys/sys_schedule.c` file where you perform the `sys_schedule()` function we put in the free field `m9_15` of the message in  $\text{fss\_priority}$  that we passed through step 1 `SYS_SCHEDULE _kernel_call sys_schedule`.
- 3) Then after the call of a system the message ends in the functions `do_schedule()` and `do_schedctl()` (`SYS_XXX -> do_XXX`) (which are located in `usr/src/kernel/system` and in there we put the  $\text{fss\_priority}$  equal to the field of the message `m9_15` where we pass the  $\text{fss\_priority}$  of step 2) where these

send the fss\_priority to the sched\_process of the archive system.c (where we changed the prototype of the function sched\_proc() through the /usr/src/kernel/proto.h file so that the function sched\_proc() no longer accept 5 argumentations, where the<sup>5th</sup> argument will be the fss\_priority) and there each process is useful he's going through the fss\_priority by updating the board. percentage h.

Having now for each user process the fss\_priority (this is ensured through the schedule\_process() function of the usr/src/servers/schedul/schedule.c, where if we have a user process then we call once to send to the core the fss\_priority of all the user processes) so that the core can choose to perform the user process with the smallest fss\_priority.

Then to select the kernel process with the lowest fss\_priority we changed the pick\_proc() function to /usr/src/kernel/proc.c. First we had before for each tail the 1st process of each tail that was ready to be executed, we added and have the last process of each tail that is ready to be executed(rdy\_tail).

η Method does what it did before. If we find out we're in the then we check if the 1st process of the user's tail is non-current (i.e. it is ready for execution), if it is not then we continue to the next tail (essentially we leave the time) while if it is ready for execution then the min is done equal to the p\_fss\_priority of the p\_fss\_priority 1st process that is ready for execution. less than >p\_nextready 1st process and if it is then we will perform some of these processes.

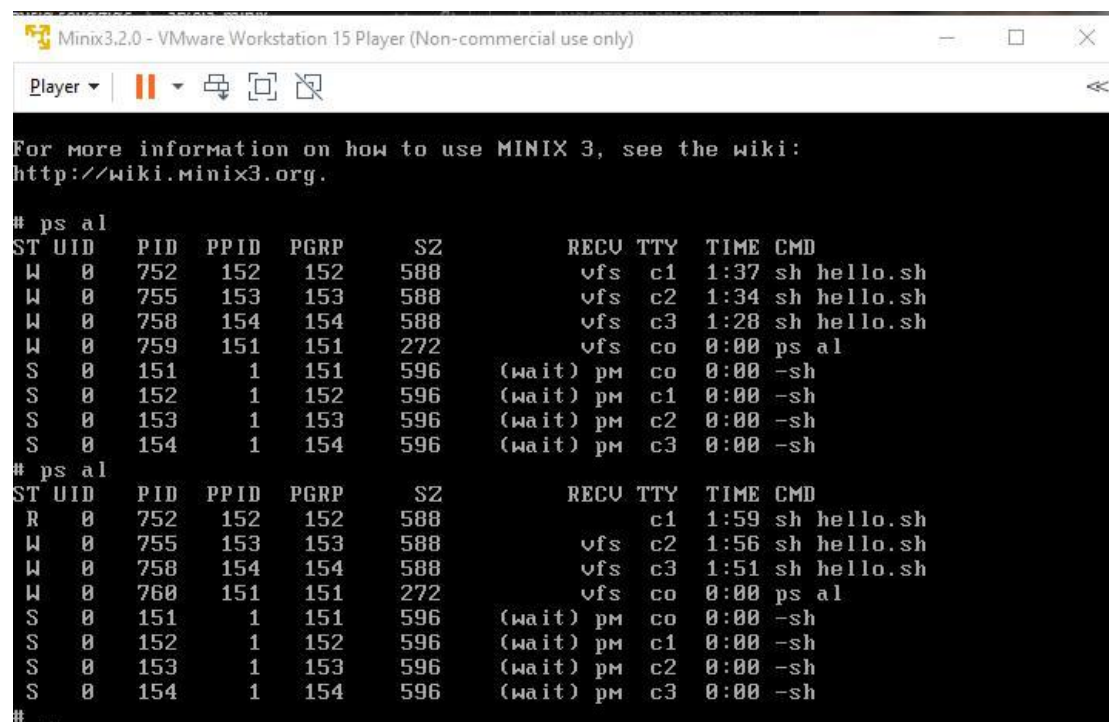
will be ready for execution). First the p will be equal to the<sup>1</sup> the process of the tail that is ready to be performed(rdy\_head[q])

while the end will be equal to the last process of the tail that is ready to be performed(i.e. rdy\_tail[q]).

## Step 4

We made 1 scrpt which we called hello.sh and ran it almost simultaneously in 3 different terminals (in the 2nd,3rd and 4th and 1st the terminal with the command ps al we saw the time of execution of the review in the terminals).

After 5 minutes (1 ps al) and 6 minutes of execution (2<sup>ps</sup> al) we had this effect here:



```
For more information on how to use MINIX 3, see the wiki:
http://wiki.minix3.org.

# ps al
ST UID  PID  PPID  PGRP  S2      RECU  TTY  TIME CMD
W  0    752   152   152   588      vfs  c1  1:37 sh hello.sh
W  0    755   153   153   588      vfs  c2  1:34 sh hello.sh
W  0    758   154   154   588      vfs  c3  1:28 sh hello.sh
W  0    759   151   151   272      vfs  co  0:00 ps al
S  0    151     1   151   596    (wait) pm  co  0:00 -sh
S  0    152     1   152   596    (wait) pm  c1  0:00 -sh
S  0    153     1   153   596    (wait) pm  c2  0:00 -sh
S  0    154     1   154   596    (wait) pm  c3  0:00 -sh

# ps al
ST UID  PID  PPID  PGRP  S2      RECU  TTY  TIME CMD
R  0    752   152   152   588      c1  1:59 sh hello.sh
W  0    755   153   153   588      vfs  c2  1:56 sh hello.sh
W  0    758   154   154   588      vfs  c3  1:51 sh hello.sh
W  0    760   151   151   272      vfs  co  0:00 ps al
S  0    151     1   151   596    (wait) pm  co  0:00 -sh
S  0    152     1   152   596    (wait) pm  c1  0:00 -sh
S  0    153     1   153   596    (wait) pm  c2  0:00 -sh
S  0    154     1   154   596    (wait) pm  c3  0:00 -sh

#
```

After 10 minutes (the last ps al shown below in the picture) we had this result here:

```

# ps al
ST UID PID PPID PGRP SZ RECU TTY TIME CMD
S 0 151 1 151 596 (wait) pm co 0:00 -sh
S 0 152 1 152 596 (wait) pm c1 0:00 -sh
S 0 153 1 153 596 (wait) pm c2 0:00 -sh
S 0 154 1 154 596 (wait) pm c3 0:00 -sh
# ps al
ST UID PID PPID PGRP SZ RECU TTY TIME CMD
W 0 752 152 152 588 vfs c1 1:59 sh hello.sh
W 0 755 153 153 588 vfs c2 1:56 sh hello.sh
W 0 758 154 154 588 vfs c3 1:51 sh hello.sh
W 0 760 151 151 272 vfs co 0:00 ps al
S 0 151 1 151 596 (wait) pm co 0:00 -sh
S 0 152 1 152 596 (wait) pm c1 0:00 -sh
S 0 153 1 153 596 (wait) pm c2 0:00 -sh
S 0 154 1 154 596 (wait) pm c3 0:00 -sh
# ps al
ST UID PID PPID PGRP SZ RECU TTY TIME CMD
W 0 752 152 152 588 vfs c1 3:39 sh hello.sh
W 0 755 153 153 588 vfs c2 3:37 sh hello.sh
W 0 758 154 154 588 vfs c3 3:33 sh hello.sh
W 0 761 151 151 272 vfs co 0:00 ps al
S 0 151 1 151 596 (wait) pm co 0:00 -sh
S 0 152 1 152 596 (wait) pm c1 0:00 -sh
S 0 153 1 153 596 (wait) pm c2 0:00 -sh
S 0 154 1 154 596 (wait) pm c3 0:00 -sh
#

```

After 15 minutes of execution we had this result here:

```

# ps al
ST UID PID PPID PGRP SZ RECU TTY TIME CMD
W 0 752 152 152 588 vfs c1 5:03 sh hello.sh
R 0 755 153 153 588 c2 5:01 sh hello.sh
R 0 758 154 154 588 c3 5:02 sh hello.sh
W 0 763 151 151 272 vfs co 0:00 ps al
S 0 151 1 151 596 (wait) pm co 0:00 -sh
S 0 152 1 152 596 (wait) pm c1 0:00 -sh
S 0 153 1 153 596 (wait) pm c2 0:00 -sh
S 0 154 1 154 596 (wait) pm c3 0:00 -sh
#

```

As shown by the above images, time is divided between the 3 different terminals that we run the script described above.